

Random forests

Jeffrey Leek, Assistant Professor of Biostatistics

May 18, 2016

Random forests

1. Bootstrap samples
2. At each split, bootstrap variables
3. Grow multiple trees and vote


Pros:

1. Accuracy

Cons:


1. Speed
2. Interpretability
3. Overfitting

Random forests



The ensemble model

Forest output probability $p(c|\mathbf{v}) = \frac{1}{T} \sum_t p_t(c|\mathbf{v})$



<http://www.robots.ox.ac.uk/~az/lectures/ml/lect5.pdf>

Iris data

```
library(caret)
```

```
## Loading required package: lattice
```

```
## Loading required package: ggplot2
```

```
data(iris); library(ggplot2)
inTrain <- createDataPartition(y=iris$Species,
                                p=0.7, list=FALSE)
training <- iris[inTrain,]
testing <- iris[-inTrain,]
```

Random forests

```
library(caret)
modFit <- train(Species~ .,data=training,method="rf",prox=7)

## Loading required package: randomForest

## randomForest 4.6-12

## Type rfNews() to see new features/changes/bug fixes.

##
## Attaching package: 'randomForest'

## The following object is masked from 'package:ggplot2':
##
##      margin

modFit
```

```
## Random Forest
```

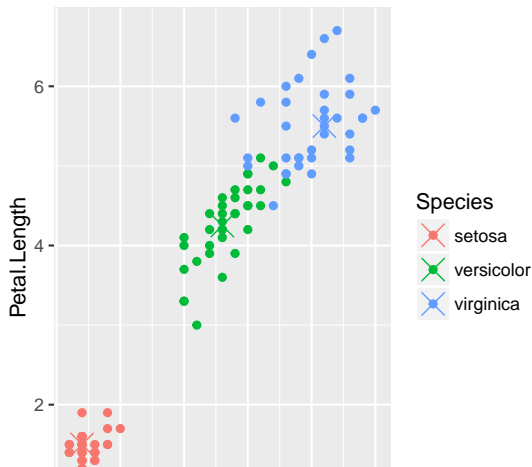
Getting a single tree

```
getTree(modFit$finalModel,k=2)
```

##	left daughter	right daughter	split var	split point	std
## 1	2	3	4	0.75	
## 2	0	0	0	0.00	
## 3	4	5	4	1.75	
## 4	6	7	3	5.35	
## 5	8	9	3	4.85	
## 6	0	0	0	0.00	
## 7	0	0	0	0.00	
## 8	0	0	0	0.00	
## 9	0	0	0	0.00	

Class “centers”

```
irisP <- classCenter(training[,c(3,4)], training$Species, n  
irisP <- as.data.frame(irisP); irisP$Species <- rownames(irisP)  
p <- qplot(Petal.Width, Petal.Length, col=Species, data=training)  
p + geom_point(aes(x=Petal.Width, y=Petal.Length, col=Species))
```




Predicting new values

```
pred <- predict(modFit,testing); testing$predRight <- pred=  
table(pred,testing$Species)
```

```
##  
## pred          setosa versicolor virginica  
## setosa         15          0          0  
## versicolor     0          15          2  
## virginica      0          0          13
```


Predicting new values

```
qplot(Petal.Width, Petal.Length, colour=predRight, data=testin
```



Notes and further resources

Notes:

- ▶ Random forests are usually one of the two top performing algorithms along with boosting in prediction contests.
- ▶ Random forests are difficult to interpret but often very accurate.
- ▶ Care should be taken to avoid overfitting (see `rfcv` function)

Further resources:

- ▶ Random forests
- ▶ Random forest Wikipedia
- ▶ Elements of Statistical Learning