

Introduction to the R Language

Roger D. Peng, Associate Professor of Biostatistics

May 17, 2016

Looping on the Command Line

Writing for, while loops is useful when programming but not particularly easy when working interactively on the command line. There are some functions which implement looping to make life easier.

- ▶ `lapply`: Loop over a list and evaluate a function on each element
- ▶ `sapply`: Same as `lapply` but try to simplify the result
- ▶ `apply`: Apply a function over the margins of an array
- ▶ `tapply`: Apply a function over subsets of a vector
- ▶ `mapply`: Multivariate version of `lapply`

An auxiliary function `split` is also useful, particularly in conjunction with `lapply`.

lapply

`lapply` takes three arguments: (1) a list `X`; (2) a function (or the name of a function) `FUN`; (3) other arguments via its `...` argument. If `X` is not a list, it will be coerced to a list using `as.list`.

```
lapply
```

```
## function (X, FUN, ...)
## {
##     FUN <- match.fun(FUN)
##     if (!is.vector(X) || is.object(X))
##         X <- as.list(X)
##     .Internal(lapply(X, FUN))
## }
## <bytecode: 0x7fcd23943780>
## <environment: namespace:base>
```

The actual looping is done internally in C code.

lapply

`lapply` always returns a list, regardless of the class of the input.

```
x <- list(a = 1:5, b = rnorm(10))  
lapply(x, mean)
```

```
## $a  
## [1] 3  
##  
## $b  
## [1] -0.02441855
```

lapply

```
x <- list(a = 1:4, b = rnorm(10), c = rnorm(20, 1), d = rnorm(10, 5))  
lapply(x, mean)
```

```
## $a  
## [1] 2.5  
##  
## $b  
## [1] -0.4948272  
##  
## $c  
## [1] 1.042135  
##  
## $d  
## [1] 5.16813
```

lapply

```
> x <- 1:4  
> lapply(x, runif)  
[[1]]  
[1] 0.2675082  
  
[[2]]  
[1] 0.2186453 0.5167968  
  
[[3]]  
[1] 0.2689506 0.1811683 0.5185761  
  
[[4]]  
[1] 0.5627829 0.1291569 0.2563676 0.7179353
```

lapply

```
> x <- 1:4  
> lapply(x, runif, min = 0, max = 10)  
[[1]]  
[1] 3.302142  
  
[[2]]  
[1] 6.848960 7.195282  
  
[[3]]  
[1] 3.5031416 0.8465707 9.7421014  
  
[[4]]  
[1] 1.195114 3.594027 2.930794 2.766946
```

lapply

`lapply` and friends make heavy use of *anonymous* functions.

```
> x <- list(a = matrix(1:4, 2, 2), b = matrix(1:6, 3, 2))
> x
$a
      [,1] [,2]
[1,]     1     3
[2,]     2     4

$b
      [,1] [,2]
[1,]     1     4
[2,]     2     5
[3,]     3     6
```


lapply

An anonymous function for extracting the first column of each matrix.

```
> lapply(x, function(elt) elt[,1])  
$a  
[1] 1 2  
  
$b  
[1] 1 2 3
```

sapply

sapply will try to simplify the result of lapply if possible.

- ▶ If the result is a list where every element is length 1, then a vector is returned
- ▶ If the result is a list where every element is a vector of the same length (> 1), a matrix is returned.
- ▶ If it can't figure things out, a list is returned

apply

```
> x <- list(a = 1:4, b = rnorm(10), c = rnorm(20, 1), d = rnorm(10, 5))
> lapply(x, mean)
$a
[1] 2.5

$b
[1] 0.06082667

$c
[1] 1.467083

$d
[1] 5.074749
```

sapply

```
> sapply(x, mean)
```

a	b	c	d
2.50000000	0.06082667	1.46708277	5.07474950

```
> mean(x)
```

```
[1] NA
```

Warning message:

In `mean.default(x)` : argument is not numeric or logical: re