

# Structure of a Data Analysis

Roger D. Peng, Associate Professor of Biostatistics

May 18, 2016

# Steps in a data analysis

- ▶ Define the question
- ▶ Define the ideal data set
- ▶ Determine what data you can access
- ▶ Obtain the data
- ▶ Clean the data
- ▶ Exploratory data analysis
- ▶ Statistical prediction/modeling
- ▶ Interpret results
- ▶ Challenge results
- ▶ Synthesize/write up results
- ▶ Create reproducible code

# Steps in a data analysis

- ▶ Define the question
- ▶ Define the ideal data set
- ▶ Determine what data you can access
- ▶ Obtain the data
- ▶ Clean the data
- ▶ Exploratory data analysis
- ▶ Statistical prediction/modeling
- ▶ Interpret results
- ▶ Challenge results
- ▶ Synthesize/write up results
- ▶ Create reproducible code

# An example


## **Start with a general question**

Can I automatically detect emails that are SPAM or not?

## **Make it concrete**

Can I use quantitative characteristics of the emails to classify them as SPAM/HAM?

# Our data set



## Spam E-mail Database

### Description

A data set collected at Hewlett-Packard Labs, that classifies 4601 e-mails as spam or non-spam. In addition to this class label there are 57 variables indicating the frequency of certain words and characters in the e-mail.

### Usage

```
data(spam)
```

### Format

A data frame with 4601 observations and 58 variables.

The first 48 variables contain the frequency of the variable name (e.g., business) in the e-mail. If the variable name starts with num (e.g., num650) the it indicates the frequency of the corresponding number (e.g., 650). The variables 49-54 indicate the frequency of the characters `.', `(' , `)` , `!`, `S`, and `#`. The variables 55-57 contain the average, longest and total run-length of capital letters. Variable 58 indicates the type of the mail and is either "nonspam" or "spam", i.e. unsolicited commercial e-mail.

### Details

The data set contains 2788 e-mails classified as "nonspam" and 1813 classified as "spam".

The "spam" concept is diverse: advertisements for products/web sites, make money fast schemes, chain letters, pornography... This collection of spam e-mails came from the collectors' postmaster and individuals who had filed spam. The collection of non-spam e-mails came from filed work and personal e-mails, and hence the word 'george' and the area code '650' are indicators of non-spam. These are useful when constructing a personalized spam filter. One would either have to blind such non-spam indicators or get a very wide collection of non-spam to generate a general purpose spam filter.

### Source

- Creators: Mark Hopkins, Erik Reeber, George Forman, Jaap Suermondt at Hewlett-Packard Labs, 1501 Page Mill Rd., Palo Alto, CA 94304
- Donor: George Forman (gforman at nospam hpl.hp.com) 650-857-7835

These data have been taken from the UCI Repository Of Machine Learning Databases at <http://www.ics.uci.edu/~mllearn/MLRepository.html>

### References

T. Hastie, R. Tibshirani, J.H. Friedman. *The Elements of Statistical Learning*. Springer, 2001.

http:  
//search.r-project.org/library/kernlab/html/spam.html

## Subsampling our data set

We need to generate a test and training set (prediction)

```
# If it isn't installed, install the kernlab package  
library(kernlab)  
data(spam)  
# Perform the subsampling  
set.seed(3435)  
trainIndicator = rbinom(4601,size=1,prob=0.5)  
table(trainIndicator)
```

```
## trainIndicator  
##      0      1  
## 2314 2287
```

```
trainSpam = spam[trainIndicator==1,]  
testSpam = spam[trainIndicator==0,]
```

# Exploratory data analysis

- ▶ Look at summaries of the data
- ▶ Check for missing data
- ▶ Create exploratory plots
- ▶ Perform exploratory analyses (e.g. clustering)

# Names

```
names(trainSpam)
```

```
## [1] "make"          "address"       "all"
## [4] "num3d"         "our"           "over"
## [7] "remove"       "internet"      "order"
## [10] "mail"         "receive"       "will"
## [13] "people"       "report"        "addresses"
## [16] "free"         "business"      "email"
## [19] "you"          "credit"        "your"
## [22] "font"         "num000"        "money"
## [25] "hp"           "hpl"           "george"
## [28] "num650"       "lab"           "labs"
## [31] "telnet"       "num857"        "data"
## [34] "num415"       "num85"         "technology"
## [37] "num1999"      "parts"         "pm"
## [40] "direct"       "cs"            "meeting"
## [43] "original"     "project"       "re"
## [46] "edu"          "table"         "conference"
```



# Head

```
head(trainSpam)
```

```
##      make address  all num3d  our over remove internet orc
## 1  0.00      0.64 0.64      0 0.32 0.00   0.00          0 0
## 7  0.00      0.00 0.00      0 1.92 0.00   0.00          0 0
## 9  0.15      0.00 0.46      0 0.61 0.00   0.30          0 0
## 12 0.00      0.00 0.25      0 0.38 0.25   0.25          0 0
## 14 0.00      0.00 0.00      0 0.90 0.00   0.90          0 0
## 16 0.00      0.42 0.42      0 1.27 0.00   0.42          0 0
##      will people report addresses free business email  you
## 1  0.64      0.00      0          0 0.32          0  1.29 1.93
## 7  1.28      0.00      0          0 0.96          0  0.32 3.85
## 9  0.92      0.00      0          0 0.00          0  0.15 1.23
## 12 0.12      0.12      0          0 0.00          0  0.00 1.16
## 14 0.00      0.90      0          0 0.00          0  0.00 2.72
## 16 0.00      0.00      0          0 1.27          0  0.00 1.70
##      num000 money hp hpl  george num650 lab labs telnet num
## 1          0  0.00  0  0          0          0          0          0
```


# Summaries

```
table(trainSpam$type)
```

```
##  
## nonspam    spam  
##    1381    906
```

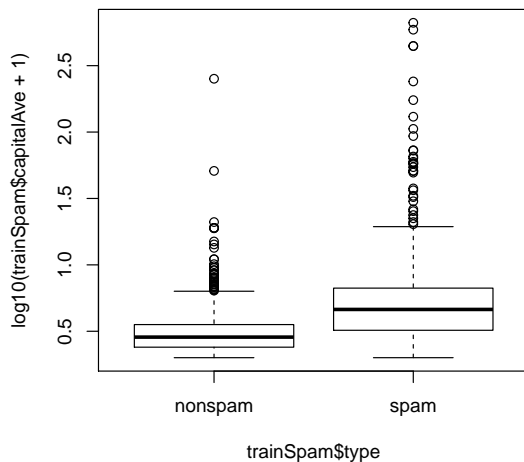
## Plots

```
plot(trainSpam$capitalAve ~ trainSpam$type)
```




# Plots

```
plot(log10(trainSpam$capitalAve + 1) ~ trainSpam$type)
```



# Relationships between predictors


```
plot(log10(trainSpam[,1:4]+1))
```



# Clustering

```
hCluster = hclust(dist(t(trainSpam[,1:57])))  
plot(hCluster)
```


Cluster Dendrogram



# New clustering

```
hClusterUpdated = hclust(dist(t(log10(trainSpam[,1:55]+1))))  
plot(hClusterUpdated)
```

Cluster Dendrogram



# Statistical prediction/modeling

- ▶ Should be informed by the results of your exploratory analysis
- ▶ Exact methods depend on the question of interest
- ▶ Transformations/processing should be accounted for when necessary
- ▶ Measures of uncertainty should be reported



# Statistical prediction/modeling

```
trainSpam$numType = as.numeric(trainSpam$type)-1
costFunction = function(x,y) sum(x!=(y > 0.5))
cvError = rep(NA,55)
library(boot)
for(i in 1:55){
  lmFormula = reformulate(names(trainSpam)[i], response = '
  glmFit = glm(lmFormula,family="binomial",data=trainSpam)
  cvError[i] = cv.glm(trainSpam,glmFit,costFunction,2)$delta
}

## Which predictor has minimum cross-validated error?
names(trainSpam)[which.min(cvError)]

## [1] "charDollar"
```

## Get a measure of uncertainty

```
## Use the best model from the group
predictionModel = glm(numType ~ charDollar,family="binomial")

## Get predictions on the test set
predictionTest = predict(predictionModel,testSpam)
predictedSpam = rep("nonspam",dim(testSpam)[1])

## Classify as `spam' for those with prob > 0.5
predictedSpam[predictionModel$fitted > 0.5] = "spam"
```

## Get a measure of uncertainty

```
## Classification table  
table(predictedSpam, testSpam$type)
```

```
##  
## predictedSpam nonspam spam  
##      nonspam      1346   458  
##      spam         61   449
```

```
## Error rate  
(61+458)/(1346+458 + 61 + 449)
```

```
## [1] 0.2242869
```

# Interpret results

- ▶ Use the appropriate language
- ▶ describes
- ▶ correlates with/associated with
- ▶ leads to/causes
- ▶ predicts
- ▶ Give an explanation
- ▶ Interpret coefficients
- ▶ Interpret measures of uncertainty

## Our example

- ▶ The fraction of characters that are dollar signs can be used to predict if an email is Spam
- ▶ Anything with more than 6.6% dollar signs is classified as Spam
- ▶ More dollar signs always means more Spam under our prediction
- ▶ Our test set error rate was 22.4%

# Challenge results

- ▶ Challenge all steps:
- ▶ Question
- ▶ Data source
- ▶ Processing
- ▶ Analysis
- ▶ Conclusions
- ▶ Challenge measures of uncertainty
- ▶ Challenge choices of terms to include in models
- ▶ Think of potential alternative analyses

# Synthesize/write-up results

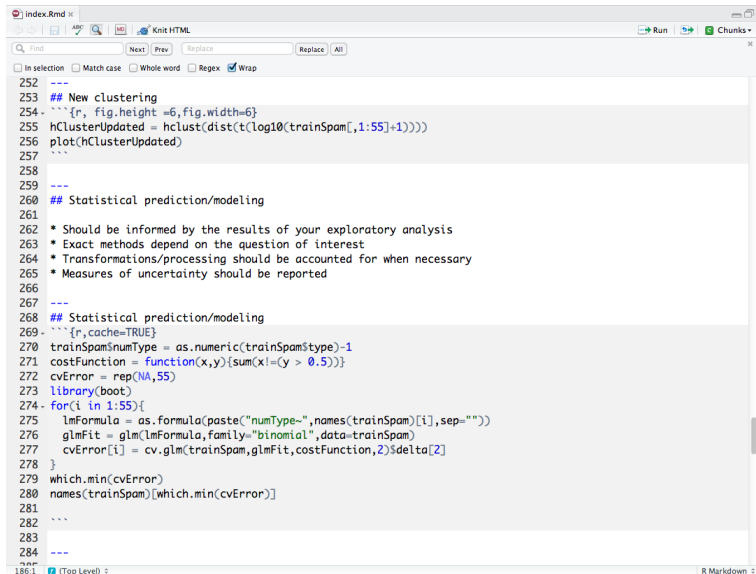
- ▶ Lead with the question
- ▶ Summarize the analyses into the story
- ▶ Don't include every analysis, include it
- ▶ If it is needed for the story
- ▶ If it is needed to address a challenge
- ▶ Order analyses according to the story, rather than chronologically
- ▶ Include “pretty” figures that contribute to the story

## In our example

- ▶ Lead with the question
- ▶ Can I use quantitative characteristics of the emails to classify them as SPAM/HAM?
- ▶ Describe the approach
- ▶ Collected data from UCI -> created training/test sets
- ▶ Explored relationships
- ▶ Choose logistic model on training set by cross validation
- ▶ Applied to test, 78% test set accuracy
- ▶ Interpret results
- ▶ Number of dollar signs seems reasonable, e.g. "Make money with Viagra \ \$ \ \$ \ \$ \ \$!"
- ▶ Challenge results
- ▶ 78% isn't that great
- ▶ I could use more variables
- ▶ Why logistic regression?



# Create reproducible code



The screenshot shows an RStudio editor window with a file named 'index.Rmd'. The editor has a search bar at the top with 'Find', 'Next', 'Prev', 'Replace', and 'All' buttons. Below the search bar are checkboxes for 'In selection', 'Match case', 'Whole word', 'Regex', and 'Wrap'. The main editor area contains R code with line numbers 252 to 285. The code is organized into sections with comments and code blocks. The first section (lines 252-257) is for clustering. The second section (lines 258-266) is for statistical prediction/modeling, with a list of bullet points. The third section (lines 267-285) is for cross-validation and model selection. The status bar at the bottom shows '186:1' and 'R Markdown'.

```
252 ---
253 ## New clustering
254 ```{r, fig.height =6,fig.width=6}
255 hClusterUpdated = hclust(dist(t(log10(trainSpam[,1:55]+1))))
256 plot(hClusterUpdated)
257 ```
258 ---
259 ---
260 ## Statistical prediction/modeling
261
262 * Should be informed by the results of your exploratory analysis
263 * Exact methods depend on the question of interest
264 * Transformations/processing should be accounted for when necessary
265 * Measures of uncertainty should be reported
266
267 ---
268 ## Statistical prediction/modeling
269 ```{r,cache=TRUE}
270 trainSpam$numType = as.numeric(trainSpam$type)-1
271 costFunction = function(x,y){sum(x!=(y > 0.5))}
272 cvError = rep(NA,55)
273 library(boot)
274 for(i in 1:55){
275   lmFormula = as.formula(paste("numType~",names(trainSpam)[i],sep=""))
276   glmFit = glm(lmFormula,family="binomial",data=trainSpam)
277   cvError[i] = cv.glm(trainSpam,glmFit,costFunction,2)$delta[2]
278 }
279 which.min(cvError)
280 names(trainSpam)[which.min(cvError)]
281
282 ```
283 ---
284 ---
285
```