

# A trip to Asymptopia

Brian Caffo, Jeff Leek, Roger Peng

May 18, 2016

# Asymptotics

- ▶ Asymptotics is the term for the behavior of statistics as the sample size (or some other relevant quantity) limits to infinity (or some other relevant number)
- ▶ (Asymptopia is my name for the land of asymptotics, where everything works out well and there's no messes. The land of infinite data is nice that way.)
- ▶ Asymptotics are incredibly useful for simple statistical inference and approximations
- ▶ (Not covered in this class) Asymptotics often lead to nice understanding of procedures
- ▶ Asymptotics generally give no assurances about finite sample performance
- ▶ The kinds of asymptotics that do are orders of magnitude more difficult to work with
- ▶ Asymptotics form the basis for frequency interpretation of probabilities (the long run proportion of times an event occurs)
- ▶ To understand asymptotics, we need a very basic understanding of limits

# Numerical limits

- ▶ Imagine a sequence
- ▶  $a_1 = .9$ ,
- ▶  $a_2 = .99$ ,
- ▶  $a_3 = .999, \dots$
- ▶ Clearly this sequence converges to 1
- ▶ Definition of a limit: For any fixed distance we can find a point in the sequence so that the sequence is closer to the limit than that distance from that point on

# Limits of random variables


- ▶ The problem is harder for random variables
- ▶ Consider  $\bar{X}_n$  the sample average of the first  $n$  of a collection of *iid* observations
- ▶ Example  $\bar{X}_n$  could be the average of the result of  $n$  coin flips (i.e. the sample proportion of heads)
- ▶ We say that  $\bar{X}_n$  converges in probability to a limit if for any fixed distance the probability of  $\bar{X}_n$  being closer (further away) than that distance from the limit converges to one (zero)

# The Law of Large Numbers

- ▶ Establishing that a random sequence converges to a limit is hard
- ▶ Fortunately, we have a theorem that does all the work for us, called the **Law of Large Numbers**
- ▶ The law of large numbers states that if  $X_1, \dots, X_n$  are iid from a population with mean  $\mu$  and variance  $\sigma^2$  then  $\bar{X}_n$  converges in probability to  $\mu$
- ▶ (There are many variations on the LLN; we are using a particularly lazy version, my favorite kind of version)

## Law of large numbers in action

```
n <- 10000; means <- cumsum(rnorm(n)) / (1 : n)
plot(1 : n, means, type = "l", lwd = 2,
     frame = FALSE, ylab = "cumulative means", xlab = "sample size",
     abline(h = 0))
```



# The Central Limit Theorem

- ▶ The **Central Limit Theorem** (CLT) is one of the most important theorems in statistics
- ▶ For our purposes, the CLT states that the distribution of averages of iid variables, properly normalized, becomes that of a standard normal as the sample size increases
- ▶ The CLT applies in an endless variety of settings
- ▶ Let  $X_1, \dots, X_n$  be a collection of iid random variables with mean  $\mu$  and variance  $\sigma^2$
- ▶ Let  $\bar{X}_n$  be their sample average
- ▶ Then  $\frac{\bar{X}_n - \mu}{\sigma/\sqrt{n}}$  has a distribution like that of a standard normal for large  $n$ .
- ▶ Remember the form

$$\frac{\bar{X}_n - \mu}{\sigma/\sqrt{n}} = \frac{\text{Estimate} - \text{Mean of estimate}}{\text{Std. Err. of estimate}}.$$

- ▶ Usually, replacing the standard error by its estimated value doesn't change the CLT




## Example

- ▶ Simulate a standard normal random variable by rolling  $n$  (six sided)
- ▶ Let  $X_i$  be the outcome for die  $i$
- ▶ Then note that  $\mu = E[X_i] = 3.5$
- ▶  $\text{Var}(X_i) = 2.92$
- ▶  $\text{SE } \sqrt{2.92/n} = 1.71/\sqrt{n}$
- ▶ Standardized mean

$$\frac{\bar{X}_n - 3.5}{1.71/\sqrt{n}}$$



# Simulation of mean of $n$ dice






# Coin CLT

- ▶ Let  $X_i$  be the 0 or 1 result of the  $i^{\text{th}}$  flip of a possibly unfair coin
- ▶ The sample proportion, say  $\hat{p}$ , is the average of the coin flips
- ▶  $E[X_i] = p$  and  $\text{Var}(X_i) = p(1 - p)$
- ▶ Standard error of the mean is  $\sqrt{p(1 - p)/n}$
- ▶ Then

$$\frac{\hat{p} - p}{\sqrt{p(1 - p)/n}}$$

will be approximately normally distributed



## CLT in practice

- ▶ In practice the CLT is mostly useful as an approximation

$$P\left(\frac{\bar{X}_n - \mu}{\sigma/\sqrt{n}} \leq z\right) \approx \Phi(z).$$

- ▶ Recall 1.96 is a good approximation to the .975<sup>th</sup> quantile of the standard normal
- ▶ Consider

$$.95 \approx P\left(-1.96 \leq \frac{\bar{X}_n - \mu}{\sigma/\sqrt{n}} \leq 1.96\right)$$

$$= P\left(\bar{X}_n + 1.96\sigma/\sqrt{n} \geq \mu \geq \bar{X}_n - 1.96\sigma/\sqrt{n}\right),$$

# Confidence intervals

- ▶ Therefore, according to the CLT, the probability that the random interval

$$\bar{X}_n \pm z_{1-\alpha/2}\sigma/\sqrt{n}$$

contains  $\mu$  is approximately  $100(1 - \alpha)\%$ , where  $z_{1-\alpha/2}$  is the  $1 - \alpha/2$  quantile of the standard normal distribution

- ▶ This is called a  $100(1 - \alpha)\%$  **confidence interval** for  $\mu$
- ▶ We can replace the unknown  $\sigma$  with  $s$

## Give a confidence interval for the average height of sons in Galton's data

```
library(UsingR); data(father.son); x <- father.son$height
```

```
## Loading required package: MASS
```

```
## Loading required package: HistData
```

```
## Loading required package: Hmisc
```

```
## Loading required package: lattice
```

```
## Loading required package: survival
```

```
## Loading required package: Formula
```

```
## Loading required package: ggplot2
```

```
##
```

```
## Attaching package: 'Hmisc'
```

## Sample proportions

- ▶ In the event that each  $X_i$  is 0 or 1 with common success probability  $p$  then  $\sigma^2 = p(1 - p)$
- ▶ The interval takes the form

$$\hat{p} \pm z_{1-\alpha/2} \sqrt{\frac{p(1-p)}{n}}$$

- ▶ Replacing  $p$  by  $\hat{p}$  in the standard error results in what is called a Wald confidence interval for  $p$
- ▶ Also note that  $p(1-p) \leq 1/4$  for  $0 \leq p \leq 1$
- ▶ Let  $\alpha = .05$  so that  $z_{1-\alpha/2} = 1.96 \approx 2$  then

$$2\sqrt{\frac{p(1-p)}{n}} \leq 2\sqrt{\frac{1}{4n}} = \frac{1}{\sqrt{n}}$$

- ▶ Therefore  $\hat{p} \pm \frac{1}{\sqrt{n}}$  is a quick CI estimate for  $p$

## Example

- ▶ Your campaign advisor told you that in a random sample of 100 likely voters, 56 intent to vote for you.
- ▶ Can you relax? Do you have this race in the bag?
- ▶ Without access to a computer or calculator, how precise is this estimate?
- ▶  $1/\sqrt{100} = .1$  so a back of the envelope calculation gives an approximate 95% interval of (0.46, 0.66)
- ▶ Not enough for you to relax, better go do more campaigning!
- ▶ Rough guidelines, 100 for 1 decimal place, 10,000 for 2, 1,000,000 for 3.

```
round(1 / sqrt(10 ^ (1 : 6)), 3)
```

```
## [1] 0.316 0.100 0.032 0.010 0.003 0.001
```

## Poisson interval

- ▶ A nuclear pump failed 5 times out of 94.32 days, give a 95% confidence interval for the failure rate per day?
- ▶  $X \sim \text{Poisson}(\lambda t)$ .
- ▶ Estimate  $\hat{\lambda} = X/t$
- ▶  $\text{Var}(\hat{\lambda}) = \lambda/t$

$$\frac{\hat{\lambda} - \lambda}{\sqrt{\hat{\lambda}/t}} = \frac{X - t\lambda}{\sqrt{X}} \rightarrow N(0, 1)$$

- ▶ This isn't the best interval.
- ▶ There are better asymptotic intervals.
- ▶ You can get an exact CI in this case.

## R code

```
x <- 5; t <- 94.32; lambda <- x / t  
round(lambda + c(-1, 1) * qnorm(.975) * sqrt(lambda / t), 3)
```



## In the regression class

```
exp(confint(glm(x ~ 1 + offset(log(t)), family = poisson(1
```

```
## Waiting for profiling to be done...
```

```
##      2.5 %      97.5 %
```

```
## 0.01900677 0.11393446
```