

Likelihood

Brian Caffo, Roger Peng, Jeff Leek

May 18, 2016

Likelihood

- ▶ A common and fruitful approach to statistics is to assume that the data arises from a family of distributions indexed by a parameter that represents a useful summary of the distribution
- ▶ The **likelihood** of a collection of data is the joint density evaluated as a function of the parameters with the data fixed
- ▶ Likelihood analysis of data uses the likelihood to perform inference regarding the unknown parameter

Likelihood

Given a statistical probability mass function or density, say $f(x, \theta)$, where θ is an unknown parameter, the **likelihood** is f viewed as a function of θ for a fixed, observed value of x .

Interpretations of likelihoods

The likelihood has the following properties:

1. Ratios of likelihood values measure the relative evidence of one value of the unknown parameter to another.
2. Given a statistical model and observed data, all of the relevant information contained in the data regarding the unknown parameter is contained in the likelihood.
3. If $\{X_i\}$ are independent random variables, then their likelihoods multiply. That is, the likelihood of the parameters given all of the X_i is simply the product of the individual likelihoods.

Example

- ▶ Suppose that we flip a coin with success probability θ
- ▶ Recall that the mass function for x

$$f(x, \theta) = \theta^x (1 - \theta)^{1-x} \quad \text{for } \theta \in [0, 1].$$

where x is either 0 (Tails) or 1 (Heads)

- ▶ Suppose that the result is a head
- ▶ The likelihood is

$$\mathcal{L}(\theta, 1) = \theta^1 (1 - \theta)^{1-1} = \theta \quad \text{for } \theta \in [0, 1].$$

- ▶ Therefore, $\mathcal{L}(.5, 1)/\mathcal{L}(.25, 1) = 2$,
- ▶ There is twice as much evidence supporting the hypothesis that $\theta = .5$ to the hypothesis that $\theta = .25$

Example continued

- ▶ Suppose now that we flip our coin from the previous example 4 times and get the sequence 1, 0, 1, 1
- ▶ The likelihood is:


$$\begin{aligned}\mathcal{L}(\theta, 1, 0, 1, 1) &= \theta^1(1 - \theta)^{1-1}\theta^0(1 - \theta)^{1-0} \\ &\times \theta^1(1 - \theta)^{1-1}\theta^1(1 - \theta)^{1-1} \\ &= \theta^3(1 - \theta)^1\end{aligned}$$

- ▶ This likelihood only depends on the total number of heads and the total number of tails; we might write $\mathcal{L}(\theta, 1, 3)$ for shorthand
- ▶ Now consider $\mathcal{L}(.5, 1, 3)/\mathcal{L}(.25, 1, 3) = 5.33$
- ▶ There is over five times as much evidence supporting the hypothesis that $\theta = .5$ over that $\theta = .25$

Plotting likelihoods

- ▶ Generally, we want to consider all the values of θ between 0 and 1
- ▶ A **likelihood plot** displays θ by $\mathcal{L}(\theta, x)$
- ▶ Because the likelihood measures *relative evidence*, dividing the curve by its maximum value (or any other value for that matter) does not change its interpretation

```
pvals <- seq(0, 1, length = 1000)
plot(pvals, dbinom(3, 4, pvals) / dbinom(3, 4, 3/4), type =
```



Maximum likelihood

- ▶ The value of θ where the curve reaches its maximum has a special meaning
- ▶ It is the value of θ that is most well supported by the data
- ▶ This point is called the **maximum likelihood estimate** (or MLE) of θ

$$MLE = \operatorname{argmax}_{\theta} \mathcal{L}(\theta, x).$$

- ▶ Another interpretation of the MLE is that it is the value of θ that would make the data that we observed most probable

Some results

- ▶ $X_1, \dots, X_n \stackrel{iid}{\sim} N(\mu, \sigma^2)$ the MLE of μ is \bar{X} and the ML of σ^2 is the biased sample variance estimate.
- ▶ If $X_1, \dots, X_n \stackrel{iid}{\sim} \text{Bernoulli}(p)$ then the MLE of p is \bar{X} (the sample proportion of 1s).
- ▶ If $X_i \stackrel{iid}{\sim} \text{Binomial}(n_i, p)$ then the MLE of p is $\frac{\sum_{i=1}^n X_i}{\sum_{i=1}^n n_i}$ (the sample proportion of 1s).
- ▶ If $X \stackrel{iid}{\sim} \text{Poisson}(\lambda t)$ then the MLE of λ is X/t .
- ▶ If $X_i \stackrel{iid}{\sim} \text{Poisson}(\lambda t_i)$ then the MLE of λ is $\frac{\sum_{i=1}^n X_i}{\sum_{i=1}^n t_i}$

Example

- ▶ You saw 5 failure events per 94 days of monitoring a nuclear pump.
- ▶ Assuming Poisson, plot the likelihood

```
lambda <- seq(0, .2, length = 1000)
likelihood <- dpois(5, 94 * lambda) / dpois(5, 5)
plot(lambda, likelihood, frame = FALSE, lwd = 3, type = "l")
lines(rep(5/94, 2), 0 : 1, col = "red", lwd = 3)
lines(range(lambda[likelihood > 1/16]), rep(1/16, 2), lwd = 3)
lines(range(lambda[likelihood > 1/8]), rep(1/8, 2), lwd = 3)
```

