

# Multiple testing

Brian Caffo, Jeffrey Leek, Roger Peng

May 19, 2016

# Key ideas

- ▶ Hypothesis testing/significance analysis is commonly overused
- ▶ Correcting for multiple testing avoids false positives or discoveries
- ▶ Two key components
  - ▶ Error measure
  - ▶ Correction

## Three eras of statistics

**The age of Quetelet and his successors, in which huge census-level data sets were brought to bear on simple but important questions:** Are there more male than female births? Is the rate of insanity rising?


The classical period of Pearson, Fisher, Neyman, Hotelling, and their successors, intellectual giants who **developed a theory of optimal inference capable of wringing every drop of information out of a scientific experiment.** The questions dealt with still tended to be simple Is treatment A better than treatment B?

**The era of scientific mass production,** in which new technologies typified by the microarray allow a single team of scientists to produce data sets of a size Quetelet would envy. But now the flood of data is accompanied by a deluge of questions, perhaps thousands of estimates or hypothesis tests that the statistician is charged with answering together; not at all what the classical masters had in mind. Which variables matter among the thousands measured? How do you relate unrelated information?


# Reasons for multiple testing



# Why correct for multiple tests?



Why correct for multiple tests?



<http://xkcd.com/882/>

# Types of errors

Suppose you are testing a hypothesis that a parameter  $\beta$  equals zero versus the alternative that it does not equal zero. These are the possible outcomes.

	$\beta = 0$	$\beta \neq 0$	Hypotheses
--	-------------	----------------	------------

Claim $\beta = 0$	$U$	$T$	$m - R$
Claim $\beta \neq 0$	$V$	$S$	$R$
Claims	$m_0$		
	$m - m_0$	$m$	

**Type I error or false positive ( $V$ )** Say that the parameter does not equal zero when it does

**Type II error or false negative ( $T$ )** Say that the parameter equals zero when it doesn't

# Error rates

**False positive rate** - The rate at which false results ( $\beta = 0$ ) are called significant:  $E \left[ \frac{V}{m_0} \right]^*$

**Family wise error rate (FWER)** - The probability of at least one false positive  $\Pr(V \geq 1)$

**False discovery rate (FDR)** - The rate at which claims of significance are false  $E \left[ \frac{V}{R} \right]$

- ▶ The false positive rate is closely related to the type I error rate  
[http://en.wikipedia.org/wiki/False\\_positive\\_rate](http://en.wikipedia.org/wiki/False_positive_rate)



# Controlling the false positive rate

If P-values are correctly calculated calling all  $P < \alpha$  significant will control the false positive rate at level  $\alpha$  on average.

Problem: Suppose that you perform 10,000 tests and  $\beta = 0$  for all of them.

Suppose that you call all  $P < 0.05$  significant.

The expected number of false positives is:  $10,000 \times 0.05 = 500$  false positives.

**How do we avoid so many false positives?**

# Controlling family-wise error rate (FWER)

The Bonferroni correction is the oldest multiple testing correction.

**Basic idea:** \* Suppose you do  $m$  tests \* You want to control FWER at level  $\alpha$  so  $Pr(V \geq 1) < \alpha$  \* Calculate P-values normally \* Set  $\alpha_{fwer} = \alpha/m$  \* Call all  $P$ -values less than  $\alpha_{fwer}$  significant

**Pros:** Easy to calculate, conservative **Cons:** May be very conservative

# Controlling false discovery rate (FDR)


This is the most popular correction when performing *lots* of tests say in genomics, imaging, astronomy, or other signal-processing disciplines.

**Basic idea:** \* Suppose you do  $m$  tests \* You want to control FDR at level  $\alpha$  so  $E \left[ \frac{V}{R} \right]$  \* Calculate P-values normally \* Order the P-values from smallest to largest  $P_{(1)}, \dots, P_{(m)}$  \* Call any  $P_{(i)} \leq \alpha \times \frac{i}{m}$  significant

**Pros:** Still pretty easy to calculate, less conservative (maybe much less)

**Cons:** Allows for more false positives, may behave strangely under dependence

## Example with 10 P-values



Controlling all error rates at  $\alpha = 0.20$

# Adjusted P-values

- ▶ One approach is to adjust the threshold  $\alpha$
- ▶ A different approach is to calculate “adjusted p-values”
- ▶ They *are not p-values* anymore
- ▶ But they can be used directly without adjusting  $\alpha$

**Example:** \* Suppose P-values are  $P_1, \dots, P_m$  \* You could adjust them by taking  $P_i^{fwer} = \max(m \times P_i, 1)$  for each P-value. \* Then if you call all  $P_i^{fwer} < \alpha$  significant you will control the FWER.

## Case study I: no true positives

```
set.seed(1010093)
pValues <- rep(NA,1000)
for(i in 1:1000){
  y <- rnorm(20)
  x <- rnorm(20)
  pValues[i] <- summary(lm(y ~ x))$coeff[2,4]
}

# Controls false positive rate
sum(pValues < 0.05)
```

```
## [1] 51
```

## Case study I: no true positives

```
# Controls FWER
```

```
sum(p.adjust(pValues,method="bonferroni") < 0.05)
```

```
## [1] 0
```

```
# Controls FDR
```

```
sum(p.adjust(pValues,method="BH") < 0.05)
```

```
## [1] 0
```

## Case study II: 50% true positives

```
set.seed(1010093)
pValues <- rep(NA,1000)
for(i in 1:1000){
  x <- rnorm(20)
  # First 500 beta=0, last 500 beta=2
  if(i <= 500){y <- rnorm(20)}else{ y <- rnorm(20,mean=2*x)}
  pValues[i] <- summary(lm(y ~ x))$coeff[2,4]
}
trueStatus <- rep(c("zero","not zero"),each=500)
table(pValues < 0.05, trueStatus)
```

```
##           trueStatus
##           not zero zero
##  FALSE           0  476
##  TRUE           500   24
```



## Case study II: 50% true positives

```
# Controls FWER
```

```
table(p.adjust(pValues,method="bonferroni") < 0.05,trueStat
```

```
##           trueStatus
##           not zero zero
##  FALSE           23  500
##  TRUE            477    0
```

```
# Controls FDR
```



```
table(p.adjust(pValues,method="BH") < 0.05,trueStatus)
```

```
##           trueStatus
##           not zero zero
##  FALSE           0  487
##  TRUE           500   13
```

## Case study II: 50% true positives

### P-values versus adjusted P-values

```
par(mfrow=c(1,2))  
plot(pValues,p.adjust(pValues,method="bonferroni"),pch=19)  
plot(pValues,p.adjust(pValues,method="BH"),pch=19)
```



# Notes and resources

**Notes:** \* Multiple testing is an entire subfield \* A basic Bonferroni/BH correction is usually enough \* If there is strong dependence between tests there may be problems \* Consider method="BY"

**Further resources:** \* Multiple testing procedures with applications to genomics \* Statistical significance for genome-wide studies \* Introduction to multiple testing